

gpu code

```
%%writefile programs/meanFilter.cu
// CUDA implementation of Mean Filter
#include "opencv2/imgproc/imgproc.hpp"
#include <opencv2/highgui.hpp>
#include <iostream>
#include <string>
#include <stdio.h>
#include <cuda.h>
#include "cuda_runtime.h"

#define BLOCK_SIZE    16
#define FILTER_WIDTH  8
#define FILTER_HEIGHT 8

using namespace std;

// Run Mean Filter on GPU
__global__ void meanFilter(unsigned char *srcImage, unsigned char *dstImage, unsigned int width,
unsigned int height, int channel)
{
    int x = blockIdx.x*blockDim.x + threadIdx.x;
    int y = blockIdx.y*blockDim.y + threadIdx.y;

    // only threads inside image will write results
    if((x>=FILTER_WIDTH/2) && (x<(width-FILTER_WIDTH/2)) && (y>=FILTER_HEIGHT/2) &&
(y<(height-FILTER_HEIGHT/2)))
    {
        for(int c=0 ; c<channel ; c++)
        {
            unsigned char filterVector[FILTER_WIDTH*FILTER_HEIGHT];
            int sum = 0;
            // Loop inside the filter to sum pixel values
            for(int ky=0; ky<FILTER_HEIGHT; ky++) {
                for(int kx=0; kx<FILTER_WIDTH; kx++) {
                    filterVector[ky*FILTER_WIDTH+kx] = srcImage[((y+ky-FILTER_HEIGHT/2)*width +
(x+kx-FILTER_WIDTH/2))*channel+c];
                    sum += filterVector[ky*FILTER_WIDTH+kx];
                }
            }

            // Calculate mean value
            dstImage[(y*width+x)*channel+c] = sum / (FILTER_WIDTH*FILTER_HEIGHT);
        }
    }
}

// The wrapper to call mean filter
extern "C" void meanFilter_GPU_wrapper(const cv::Mat& input, cv::Mat& output)
{
    // Use cuda event to catch time
    cudaEvent_t start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);
```

```

// Calculate number of image channels
int channel = input.step/input.cols;

// Calculate number of input & output bytes in each block
const int inputSize = input.cols * input.rows * channel;
const int outputSize = output.cols * output.rows * channel;
unsigned char *d_input, *d_output;

// Allocate device memory
cudaMalloc<unsigned char>(&d_input,inputSize);
cudaMalloc<unsigned char>(&d_output,outputSize);

// Copy data from OpenCV input image to device memory
cudaMemcpy(d_input,input.ptr(),inputSize,cudaMemcpyHostToDevice);

// Specify block size
const dim3 block(BLOCK_SIZE,BLOCK_SIZE);

// Calculate grid size to cover the whole image
const dim3 grid((output.cols + block.x - 1)/block.x, (output.rows + block.y - 1)/block.y);

// Start time
cudaEventRecord(start);

// Run Mean Filter kernel on CUDA
meanFilter<<<grid,block>>>(d_input, d_output, output.cols, output.rows, channel);

// Stop time
cudaEventRecord(stop);

//Copy data from device memory to output image
cudaMemcpy(output.ptr(),d_output,outputSize,cudaMemcpyDeviceToHost);

//Free the device memory
cudaFree(d_input);
cudaFree(d_output);

cudaEventSynchronize(stop);
float milliseconds = 0;

// Calculate elapsed time in milisecond
cudaEventElapsedTime(&milliseconds, start, stop);
cout<< "\nProcessing time on GPU for mean filtering: " << milliseconds << "\n";
}

```

```

%%writefile programs/dilation.cu

```

```

#include "opencv2/imgproc/imgproc.hpp"
#include <opencv2/highgui.hpp>
#include <iostream>
#include <string>
#include <stdio.h>

```

```
#include <cuda.h>
#include "cuda_runtime.h"
```

```
#define BLOCK_SIZE    16
#define FILTER_WIDTH   8
#define FILTER_HEIGHT  8
```

```
using namespace std;
```

```
__global__ void dilationFilter(unsigned char *srcImage, unsigned char *dstImage, unsigned int width,
unsigned int height, int channel)
{
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;

    // Check if the thread is within the image boundaries
    if (x < width && y < height)
    {
        for (int c = 0; c < channel; c++)
        {
            unsigned char maxPixel = 0; // Initialize maxPixel to the minimum possible value

            // Loop through the filter window
            for (int ky = -FILTER_HEIGHT / 2; ky <= FILTER_HEIGHT / 2; ky++)
            {
                for (int kx = -FILTER_WIDTH / 2; kx <= FILTER_WIDTH / 2; kx++)
                {
                    // Calculate the neighbor's coordinates
                    int nx = x + kx;
                    int ny = y + ky;

                    // Check if the neighbor is within the image boundaries
                    if (nx >= 0 && nx < width && ny >= 0 && ny < height)
                    {
                        // Get the neighbor's pixel value
                        unsigned char neighborPixel = srcImage[(ny * width + nx) * channel + c];

                        // Update maxPixel if the neighbor's pixel value is greater
                        if (neighborPixel > maxPixel)
                        {
                            maxPixel = neighborPixel;
                        }
                    }
                }
            }

            // Set the dilation result
            dstImage[(y * width + x) * channel + c] = maxPixel;
        }
    }
}
```

```

extern "C" void dilation_GPU_wrapper(const cv::Mat& input, cv::Mat& output)
{
    // Use cuda event to catch time
    cudaEvent_t start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);

    // Calculate number of image channels
    int channel = input.step/input.cols;

    // Calculate number of input & output bytes in each block
    const int inputSize = input.cols * input.rows * channel;
    const int outputSize = output.cols * output.rows * channel;
    unsigned char *d_input, *d_output;

    // Allocate device memory
    cudaMalloc<unsigned char>(&d_input,inputSize);
    cudaMalloc<unsigned char>(&d_output,outputSize);

    // Copy data from OpenCV input image to device memory
    cudaMemcpy(d_input,input.ptr(),inputSize,cudaMemcpyHostToDevice);

    // Specify block size
    const dim3 block(BLOCK_SIZE,BLOCK_SIZE);

    // Calculate grid size to cover the whole image
    const dim3 grid((output.cols + block.x - 1)/block.x, (output.rows + block.y - 1)/block.y);

    // Start time
    cudaEventRecord(start);

    // Run Median Filter kernel on CUDA
    dilationFilter<<<grid,block>>>(d_input, d_output, output.cols, output.rows, channel);

    // Stop time
    cudaEventRecord(stop);

    //Copy data from device memory to output image
    cudaMemcpy(output.ptr(),d_output,outputSize,cudaMemcpyDeviceToHost);

    //Free the device memory
    cudaFree(d_input);
    cudaFree(d_output);

    cudaEventSynchronize(stop);
    float milliseconds = 0;

    // Calculate elapsed time in milisecond
    cudaEventElapsedTime(&milliseconds, start, stop);
    cout<< "\nProcessing time on GPU for dilation: " << milliseconds << "\n";
}

```

```

#include "opencv2/imgproc/imgproc.hpp"
#include <opencv2/highgui.hpp>
#include <iostream>
#include <string>
#include <stdio.h>
#include <cuda.h>
#include "cuda_runtime.h"
#include <cmath>

#define BLOCK_SIZE    16
#define FILTER_WIDTH  8
#define FILTER_HEIGHT 8

using namespace std;

// Run Gaussian Filter on GPU
__global__ void gaussianFilter(unsigned char *srcImage, unsigned char *dstImage, unsigned int width,
unsigned int height, int channel, float sigma) {
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y;

    __shared__ float gaussianKernel[FILTER_WIDTH * FILTER_HEIGHT];

    if (threadIdx.x < FILTER_WIDTH && threadIdx.y < FILTER_HEIGHT) {
        int index = threadIdx.y * FILTER_WIDTH + threadIdx.x;
        gaussianKernel[index] = exp(-(threadIdx.x * threadIdx.x + threadIdx.y * threadIdx.y) / (2 * sigma *
sigma)) / (2 * M_PI * sigma * sigma);
    }

    __syncthreads();

    if (x >= FILTER_WIDTH / 2 && x < width - FILTER_WIDTH / 2 && y >= FILTER_HEIGHT / 2 && y <
height - FILTER_HEIGHT / 2) {
        for (int c = 0; c < channel; c++) {
            float sum = 0.0f;
            int index = 0;
            for (int ky = -FILTER_HEIGHT / 2; ky <= FILTER_HEIGHT / 2; ky++) {
                for (int kx = -FILTER_WIDTH / 2; kx <= FILTER_WIDTH / 2; kx++) {
                    int srcX = x + kx;
                    int srcY = y + ky;
                    sum += srcImage[(srcY * width + srcX) * channel + c] * gaussianKernel[index];
                    index++;
                }
            }
            dstImage[(y * width + x) * channel + c] = sum;
        }
    }
}

```

```

// The wrapper to call Gaussian filter
extern "C" void gaussianFilter_GPU_wrapper(const cv::Mat& input, cv::Mat& output, float sigma) {
    // Use cuda event to catch time
    cudaEvent_t start, stop;
    cudaEventCreate(&start);
    cudaEventCreate(&stop);

```

```

// Calculate number of image channels
int channel = input.step / input.cols;

// Calculate number of input & output bytes in each block
const int inputSize = input.cols * input.rows * channel;
const int outputSize = output.cols * output.rows * channel;
unsigned char *d_input, *d_output;

// Allocate device memory
cudaMalloc<unsigned char>(&d_input, inputSize);
cudaMalloc<unsigned char>(&d_output, outputSize);

// Copy data from OpenCV input image to device memory
cudaMemcpy(d_input, input.ptr(), inputSize, cudaMemcpyHostToDevice);

// Specify block size
const dim3 block(BLOCK_SIZE, BLOCK_SIZE);

// Calculate grid size to cover the whole image
const dim3 grid((output.cols + block.x - 1) / block.x, (output.rows + block.y - 1) / block.y);

// Start time
cudaEventRecord(start);

// Run Gaussian Filter kernel on CUDA
gaussianFilter<<<grid, block>>>(&d_input, &d_output, output.cols, output.rows, channel, sigma);

// Stop time
cudaEventRecord(stop);

// Copy data from device memory to output image
cudaMemcpy(output.ptr(), d_output, outputSize, cudaMemcpyDeviceToHost);

// Free the device memory
cudaFree(d_input);
cudaFree(d_output);

// Synchronize the stop event
cudaEventSynchronize(stop);
float milliseconds = 0;

// Calculate elapsed time in milliseconds
cudaEventElapsedTime(&milliseconds, start, stop);
cout << "\nProcessing time on GPU for Gaussian filtering: " << milliseconds << "\n";
}

%%writefile programs/main.cpp
#include "opencv2/imgproc/imgproc.hpp"
#include <opencv2/highgui.hpp>
#include <iostream>
#include <string>
#include <stdio.h>
using namespace std;

```

```

extern "C" bool meanFilter_GPU_wrapper(const cv::Mat& input, cv::Mat& output);
extern "C" bool dilation_GPU_wrapper(const cv::Mat& input, cv::Mat& output);
extern "C" bool gaussianFilter_GPU_wrapper(const cv::Mat& input, cv::Mat& output);

// Program main
int main( int argc, char** argv ) {

    // name of image
    string image_name = "sample";
    //string image_name = "walk";

    // input & output file names
    string input_file = image_name+".jpeg";
    string output_directory = "generatedImages/";
    string output_file_meanFilter = output_directory + image_name+"_meanFilter.jpeg";
    string output_file_dilation = output_directory + image_name+"_dilation.jpeg";
    string output_file_gaussianFilter = output_directory + image_name+"_gaussianFilter.jpeg";

    // Read input image
    cv::Mat srcImage = cv::imread(input_file , cv::IMREAD_UNCHANGED);
    if(srcImage.empty())
    {
        std::cout<<"Image Not Found: "<< input_file << std::endl;
        return -1;
    }
    cout <<"\ninput image size: "<<srcImage.cols<<" "<<srcImage.rows<<" "<<srcImage.channels()<<"\n";

    // Declare the output image
    cv::Mat dstImage (srcImage.size(), srcImage.type());

    // run mean filter on GPU
    meanFilter_GPU_wrapper(srcImage, dstImage);

    // Output image after mean filter
    imwrite(output_file_meanFilter, dstImage);

    //run dilation on GPU
    dilation_GPU_wrapper(srcImage, dstImage);

    // Output image after dilation
    imwrite(output_file_dilation, dstImage);

    //run gaussianFilter on GPU
    gaussianFilter_GPU_wrapper(srcImage, dstImage);

    // Output image after gaussianFilter
    imwrite(output_file_gaussianFilter, dstImage);

    return 0;
}

```